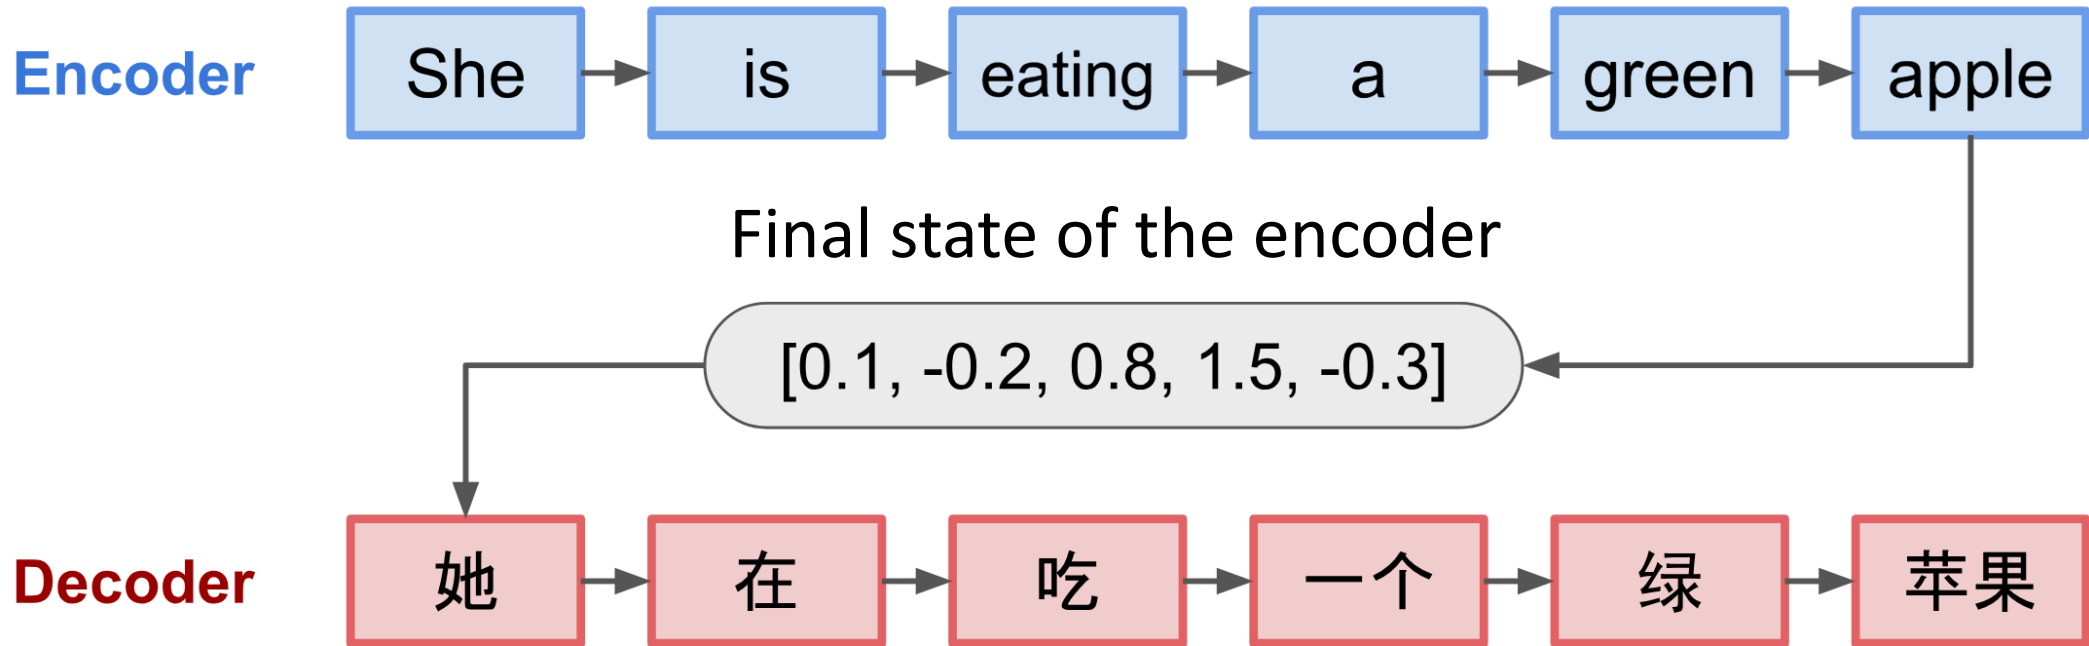


# Attention

Shusen Wang


# Seq2Seq Model



The figure is from [blog.lilianweng.github.io](http://blog.lilianweng.github.io)

# Seq2Seq Model

**Shortcoming:** The final state is incapable of remembering a **long** sequence.



The figure is from [blog.lilianweng.github.io](http://blog.lilianweng.github.io)


# Seq2Seq Model with Attention

## Original paper:

- Bahdanau, Cho, & Bengio. [Neural machine translation by jointly learning to align and translate](#). In *ICLR*, 2015.


# Seq2Seq Model

## Encoder RNN




# Seq2Seq Model

Encoder RNN



Decoder RNN




# Attention

context vector


$$\mathbf{c} = \sum_{i=1}^m \alpha_i \mathbf{h}_i.$$

- $\alpha_i$ : similarity between  $\mathbf{s}_{t-1}$  and  $\mathbf{h}_i$ .

Encoder RNN



Decoder RNN



# Attention



context vector

$$\mathbf{c} = \sum_{i=1}^m \alpha_i \mathbf{h}_i.$$

- $\alpha_i$ : similarity between  $\mathbf{s}_{t-1}$  and  $\mathbf{h}_i$ .
- $\alpha_i$  is computed by a neural network taking  $\mathbf{s}_{t-1}$  and  $\mathbf{h}_i$  as input.

**Decoder RNN**

**Encoder RNN**






# Attention

context vector

$$\mathbf{c} = \sum_{i=1}^m \alpha_i \mathbf{h}_i.$$

- $\alpha_i$ : similarity between  $\mathbf{s}_{t-1}$  and  $\mathbf{h}_i$ .
- $\alpha_i$  is computed by a neural network taking  $\mathbf{s}_{t-1}$  and  $\mathbf{h}_i$  as input.
- $\mathbf{v}$  and  $\mathbf{W}$  are trainable parameters.



# Attention

context vector

$$\mathbf{c} = \sum_{i=1}^m \alpha_i \mathbf{h}_i.$$

- $\alpha_i$ : similarity between  $\mathbf{s}_{t-1}$  and  $\mathbf{h}_i$ .





Figure is from <https://distill.pub/2016/augmented-rnns/>

# Attention


context vector

$$\mathbf{c} = \sum_{i=1}^m \alpha_i \mathbf{h}_i.$$

- $\alpha_i$ : similarity between  $\mathbf{s}_{t-1}$  and  $\mathbf{h}_i$ .




# Attention



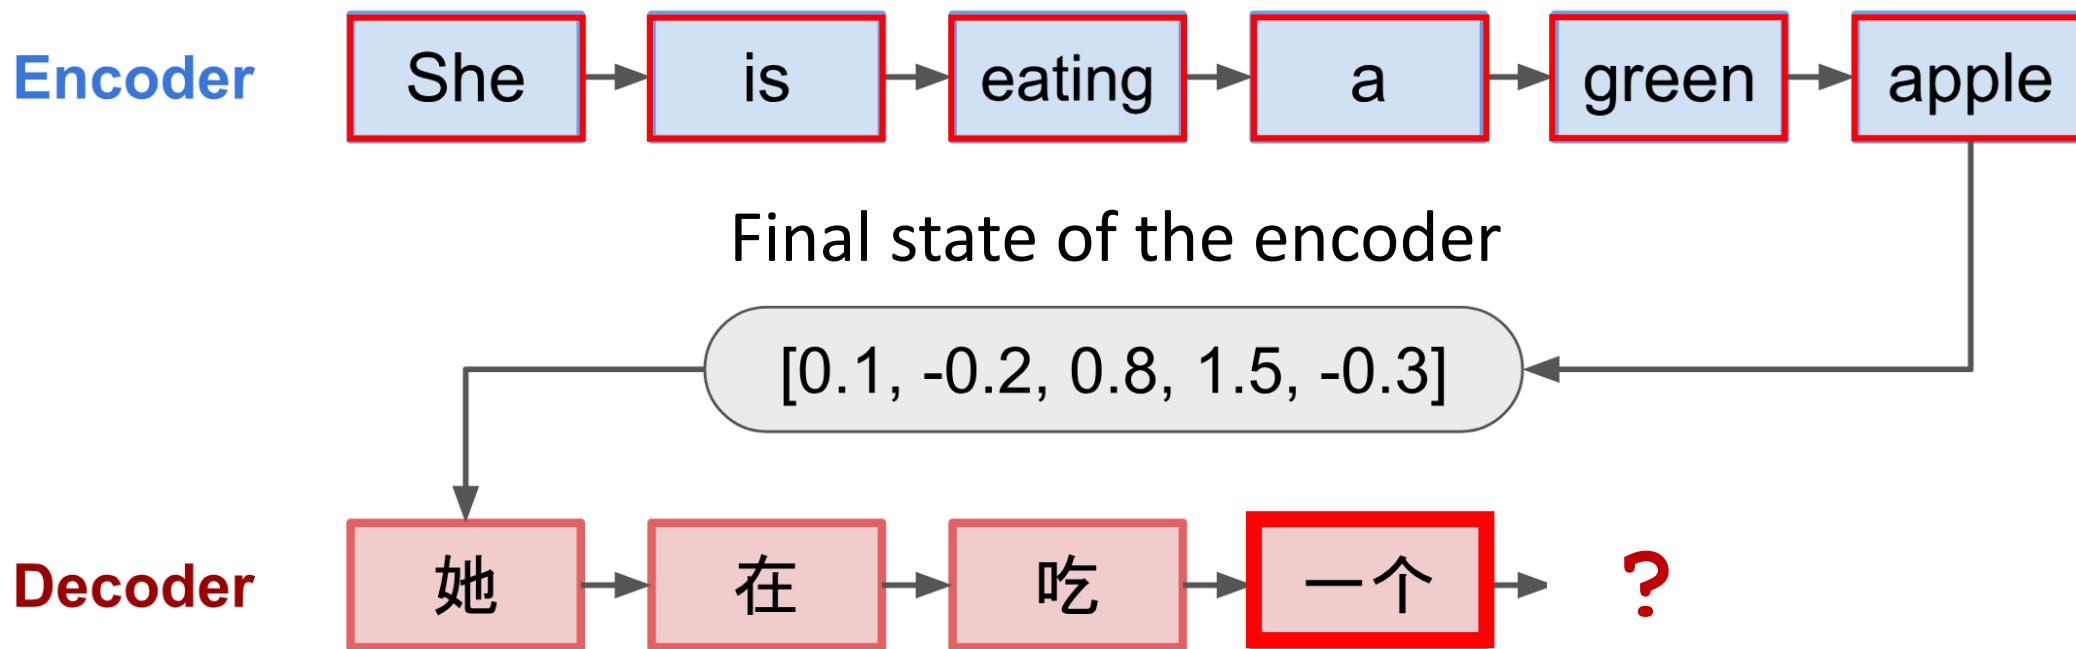
# Summary

- Standard Seq2Seq model: the decoder looks at only **its current state**.




# Summary

- Standard Seq2Seq model: the decoder looks at only its current state.
- Attention: decoder additionally looks at **all the states of the encoder**.



# Summary

- Standard Seq2Seq model: the decoder looks at only its current state.
- Attention: decoder additionally looks at **all the states of the encoder**.



# Summary

- Standard Seq2Seq model: the decoder looks at only its current state.
- Attention: decoder additionally looks at all the states of the encoder.
- Downside: higher time complexity.
  - $l_1$ : source sequence length
  - $l_2$ : target sequence length
  - Standard Seq2Seq:  $O(l_1 + l_2)$  time complexity
  - Seq2Seq + attention:  $O(l_1 l_2)$  time complexity



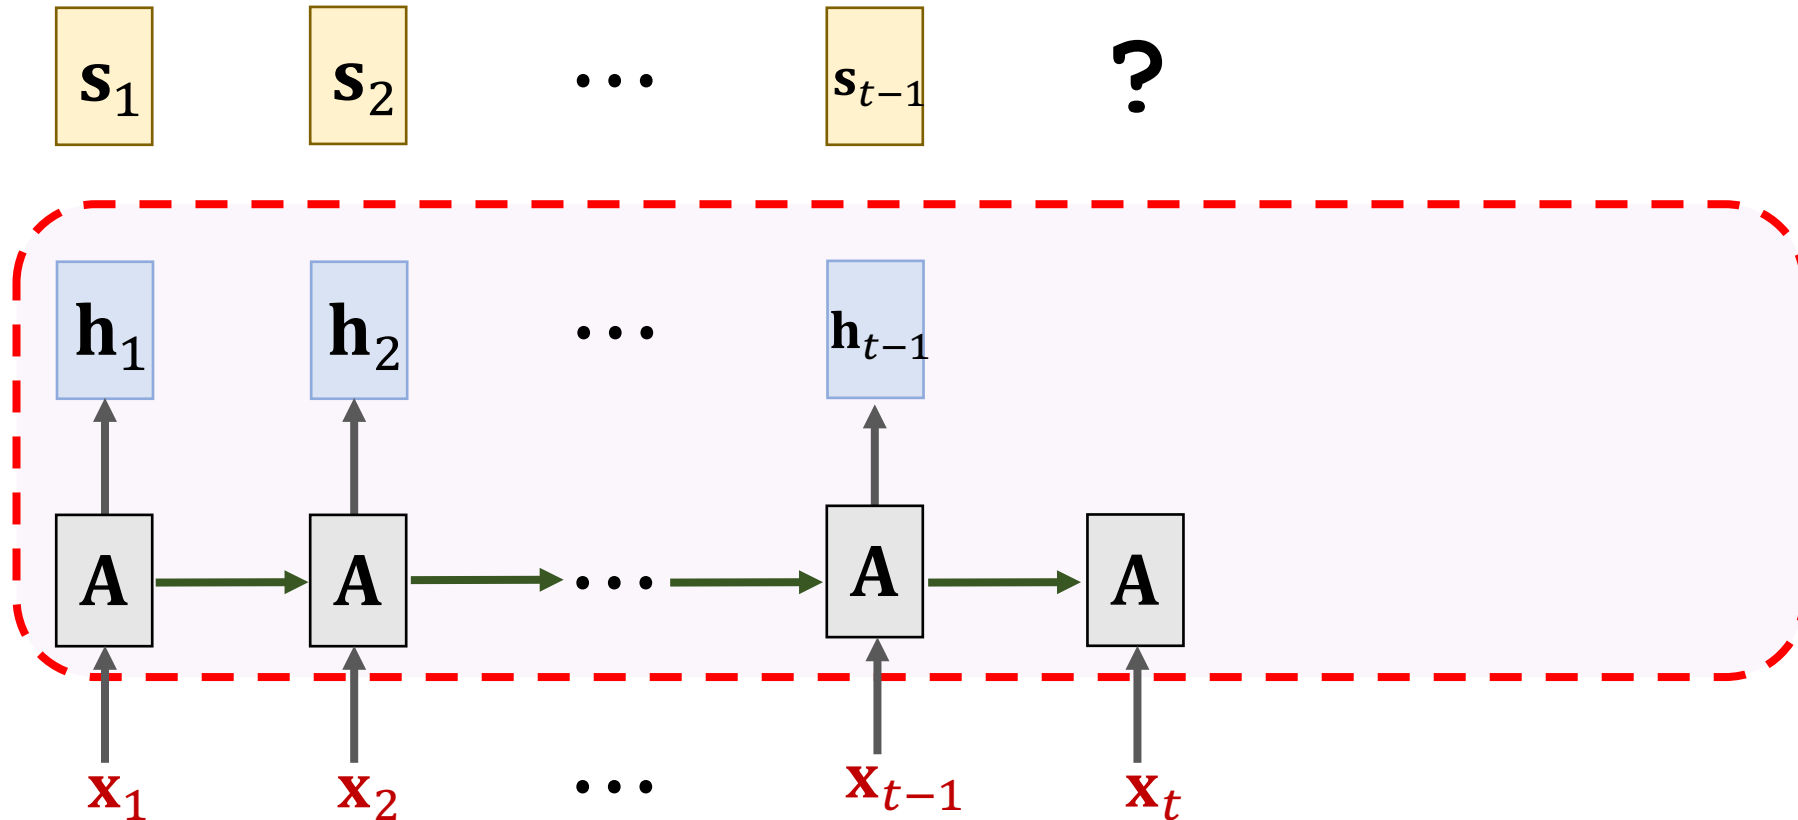
# Self-Attention: Attention beyond Seq2Seq Models

## Original paper:

- Cheng, Dong, & Lapata. [Long Short-Term Memory-Networks for Machine Reading](#). In *EMNLP*, 2016.

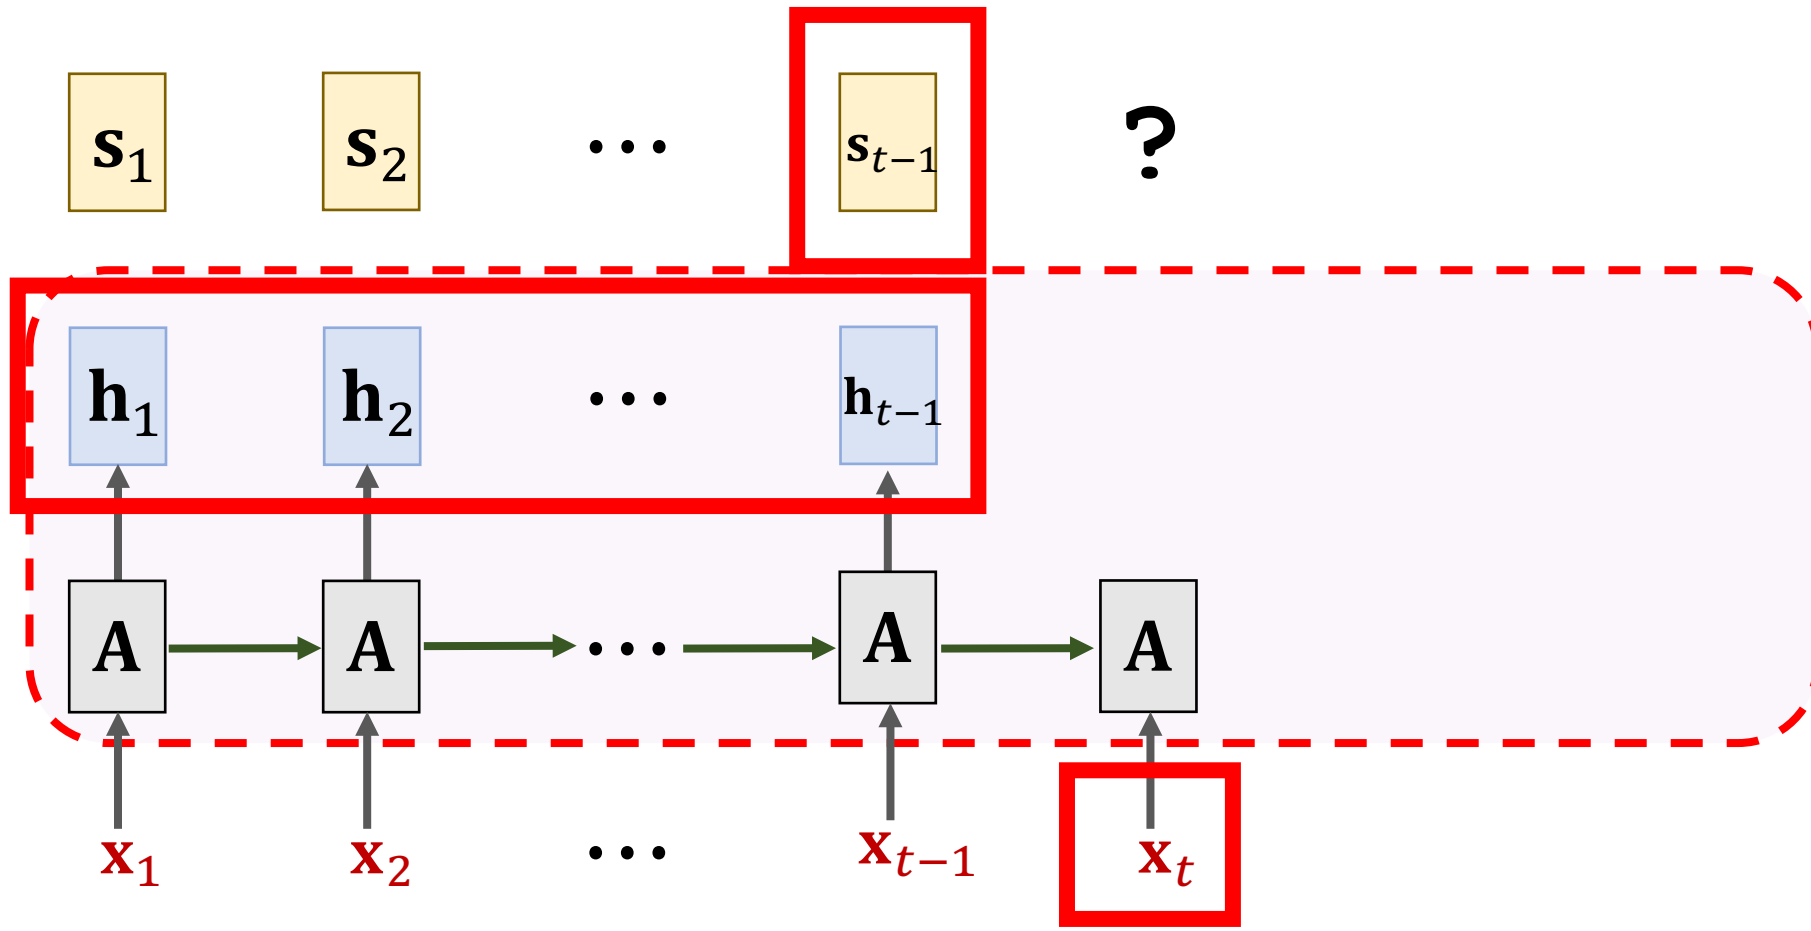
# SimpleRNN with Self-Attention

- $\mathbf{s}_t = \sum_{i=1}^{t-1} \alpha_i \mathbf{h}_i.$



# SimpleRNN with Self-Attention

- $\mathbf{s}_t = \sum_{i=1}^{t-1} \alpha_i \mathbf{h}_i$ .
- $\alpha_i$  is computed by a neural network taking  $\mathbf{x}_t$ ,  $\mathbf{s}_{t-1}$ , and  $\mathbf{h}_i$  as inputs.



# SimpleRNN with Self-Attention

- $\mathbf{s}_t = \sum_{i=1}^{t-1} \alpha_i \mathbf{h}_i$ .
- $\alpha_i$  is computed by a neural network taking  $\mathbf{x}_t$ ,  $\mathbf{s}_{t-1}$ , and  $\mathbf{h}_i$  as inputs.

$$\alpha_i = \mathbf{v}^T \cdot \tanh \left[ \mathbf{W}_x \cdot \mathbf{x}_t + \mathbf{W}_s \cdot \mathbf{s}_{t-1} + \mathbf{W}_h \cdot \mathbf{h}_i \right]$$

# SimpleRNN with Self-Attention


- $\mathbf{s}_t = \sum_{i=1}^{t-1} \alpha_i \mathbf{h}_i$ .
- $\alpha_i$  is computed by a neural network taking  $\mathbf{x}_t$ ,  $\mathbf{s}_{t-1}$ , and  $\mathbf{h}_i$  as inputs.

$$\alpha_i = \mathbf{v}^T \cdot \tanh \left[ \mathbf{W}_x \mathbf{x}_t + \mathbf{W}_s \mathbf{s}_{t-1} + \mathbf{W}_h \mathbf{h}_i \right]$$

Trainable parameters


# SimpleRNN with Self-Attention

- $\mathbf{s}_t = \sum_{i=1}^{t-1} \alpha_i \mathbf{h}_i$ .
- $\alpha_i$  is computed by a neural network taking  $\mathbf{x}_t$ ,  $\mathbf{s}_{t-1}$ , and  $\mathbf{h}_i$  as inputs.




# SimpleRNN with Self-Attention

- $\mathbf{s}_t = \sum_{i=1}^{t-1} \alpha_i \mathbf{h}_i$ .
- $\alpha_i$  is computed by a neural network taking  $\mathbf{x}_t$ ,  $\mathbf{s}_{t-1}$ , and  $\mathbf{h}_i$  as inputs.



# SimpleRNN with Self-Attention

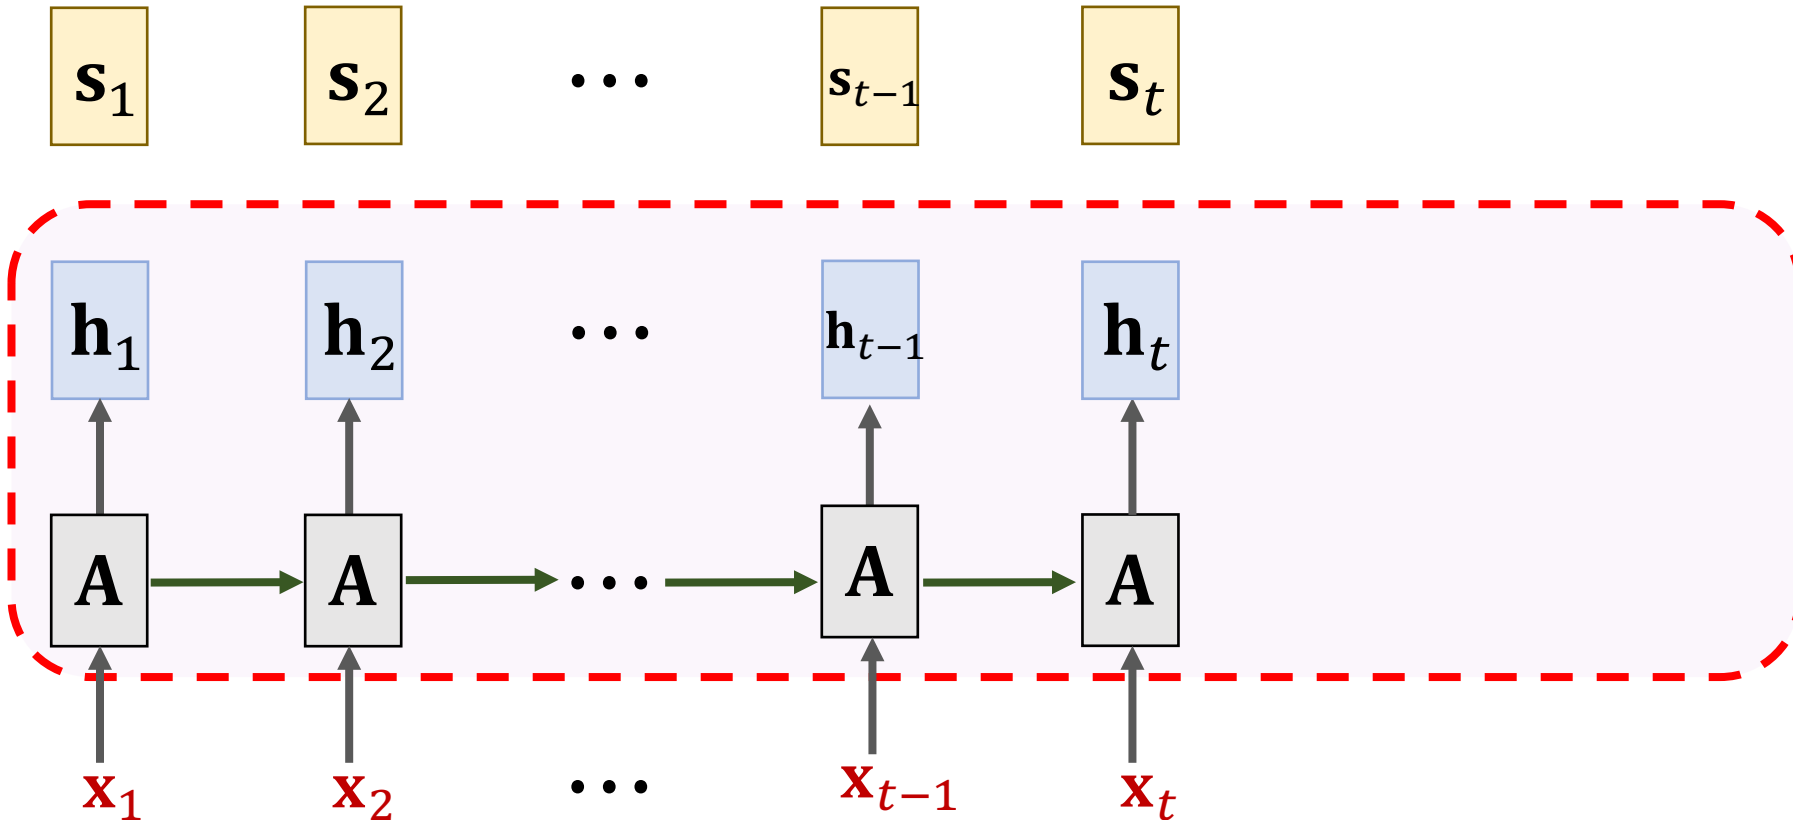
- $\mathbf{s}_t = \sum_{i=1}^{t-1} \alpha_i \mathbf{h}_i$ .
- $\alpha_i$  is computed by a neural network taking  $\mathbf{x}_t$ ,  $\mathbf{s}_{t-1}$ , and  $\mathbf{h}_i$  as inputs.





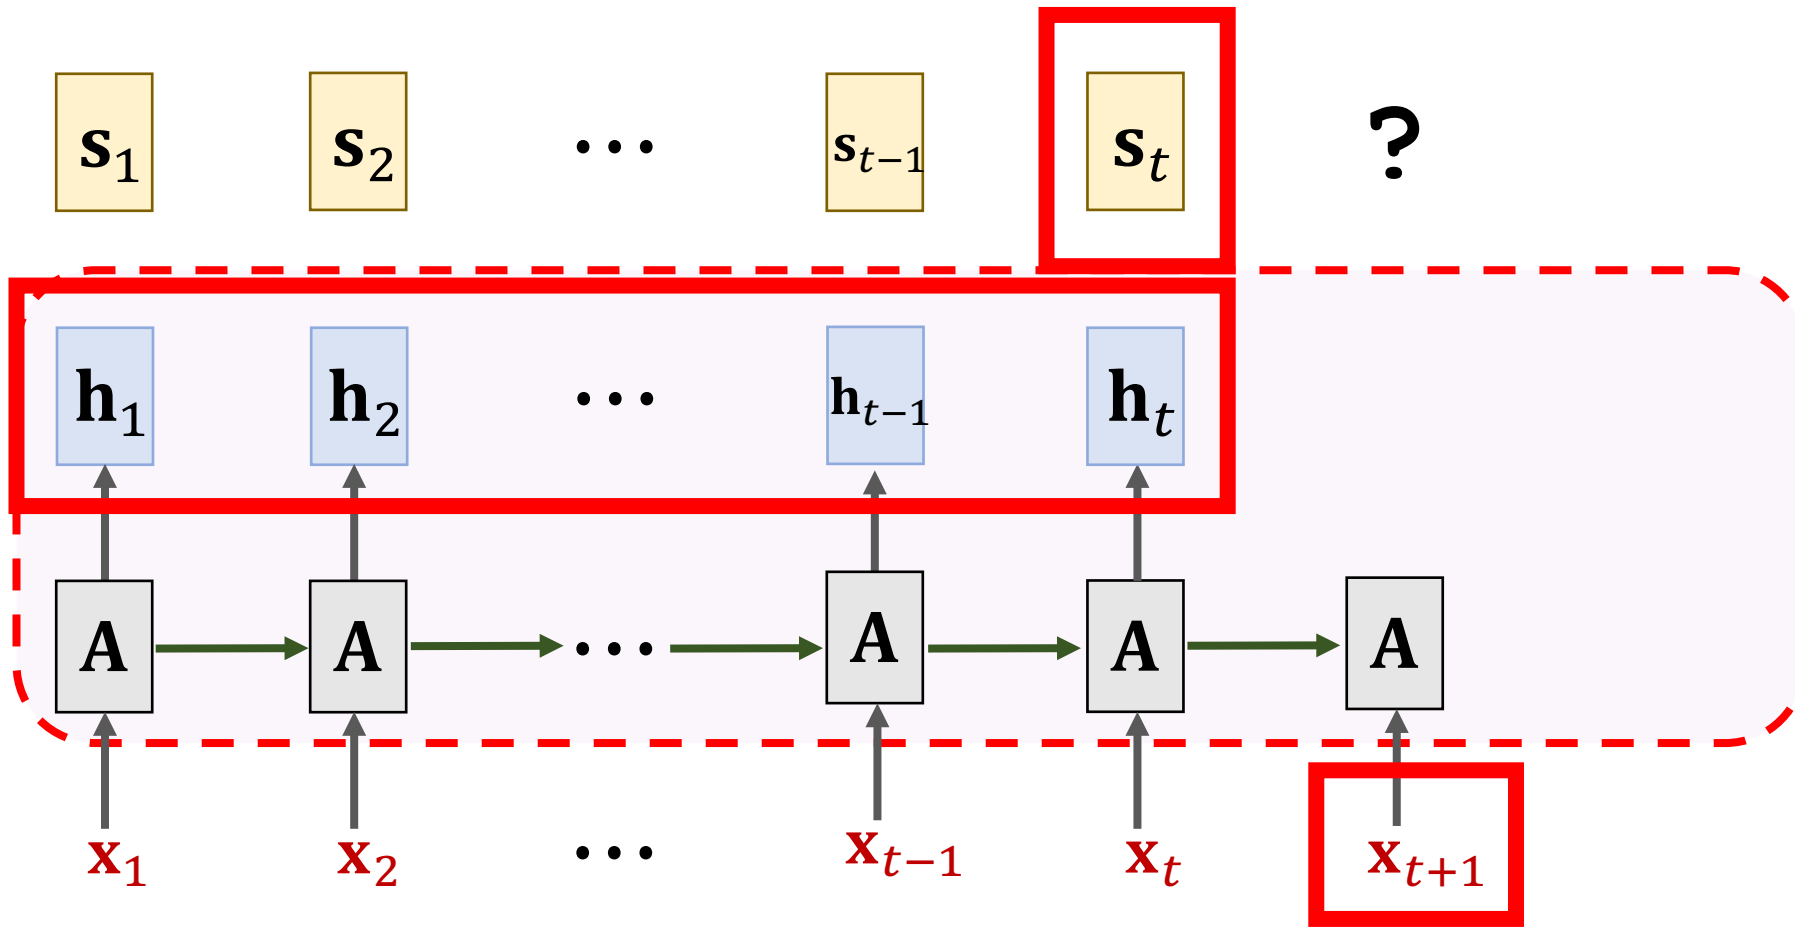
# SimpleRNN with Self-Attention

- $\mathbf{s}_t = \sum_{i=1}^{t-1} \alpha_i \mathbf{h}_i$ .
- $\alpha_i$  is computed by a neural network taking  $\mathbf{x}_t$ ,  $\mathbf{s}_{t-1}$ , and  $\mathbf{h}_i$  as inputs.



# SimpleRNN with Self-Attention

- $\mathbf{s}_{t+1} = \sum_{i=1}^t \alpha_i \mathbf{h}_i$ .
- $\alpha_i$  is computed by a neural network taking  $\mathbf{x}_{t+1}$ ,  $\mathbf{s}_t$ , and  $\mathbf{h}_i$  as inputs.



# SimpleRNN with Self-Attention


- $\mathbf{s}_{t+1} = \sum_{i=1}^t \alpha_i \mathbf{h}_i$ .
- $\alpha_i$  is computed by a neural network taking  $\mathbf{x}_{t+1}$ ,  $\mathbf{s}_t$ , and  $\mathbf{h}_i$  as inputs.

- Not the same  $\{\alpha_i\}$  computed previously.
- Because  $\mathbf{x}_t \rightarrow \mathbf{x}_{t+1}$  and  $\mathbf{s}_{t-1} \rightarrow \mathbf{s}_t$ .

$$\alpha_i = \mathbf{v}^T \cdot \tanh \left[ \mathbf{W}_x \cdot \mathbf{x}_{t+1} + \mathbf{W}_s \cdot \mathbf{s}_t + \mathbf{W}_h \cdot \mathbf{h}_i \right]$$


# SimpleRNN with Self-Attention

- $\mathbf{s}_{t+1} = \sum_{i=1}^t \alpha_i \mathbf{h}_i$ .
- $\alpha_i$  is computed by a neural network taking  $\mathbf{x}_{t+1}$ ,  $\mathbf{s}_t$ , and  $\mathbf{h}_i$  as inputs.




# SimpleRNN with Self-Attention

- $\mathbf{s}_{t+1} = \sum_{i=1}^t \alpha_i \mathbf{h}_i$ .
- $\alpha_i$  is computed by a neural network taking  $\mathbf{x}_{t+1}$ ,  $\mathbf{s}_t$ , and  $\mathbf{h}_i$  as inputs.



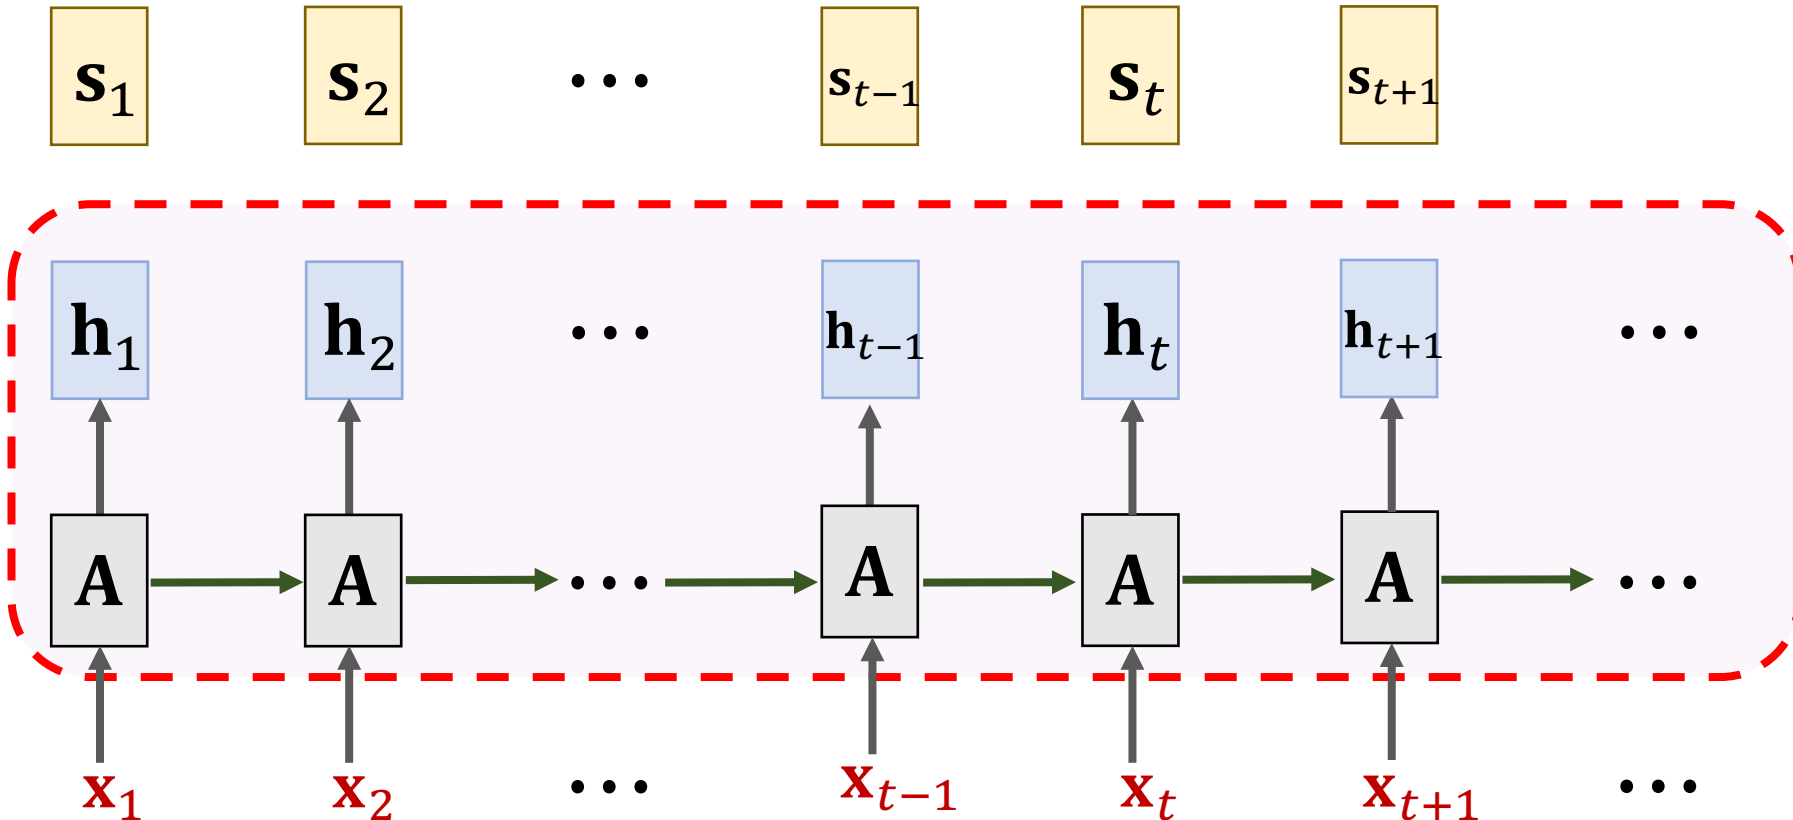
# SimpleRNN with Self-Attention

- $\mathbf{s}_{t+1} = \sum_{i=1}^t \alpha_i \mathbf{h}_i$ .
- $\alpha_i$  is computed by a neural network taking  $\mathbf{x}_{t+1}$ ,  $\mathbf{s}_t$ , and  $\mathbf{h}_i$  as inputs.

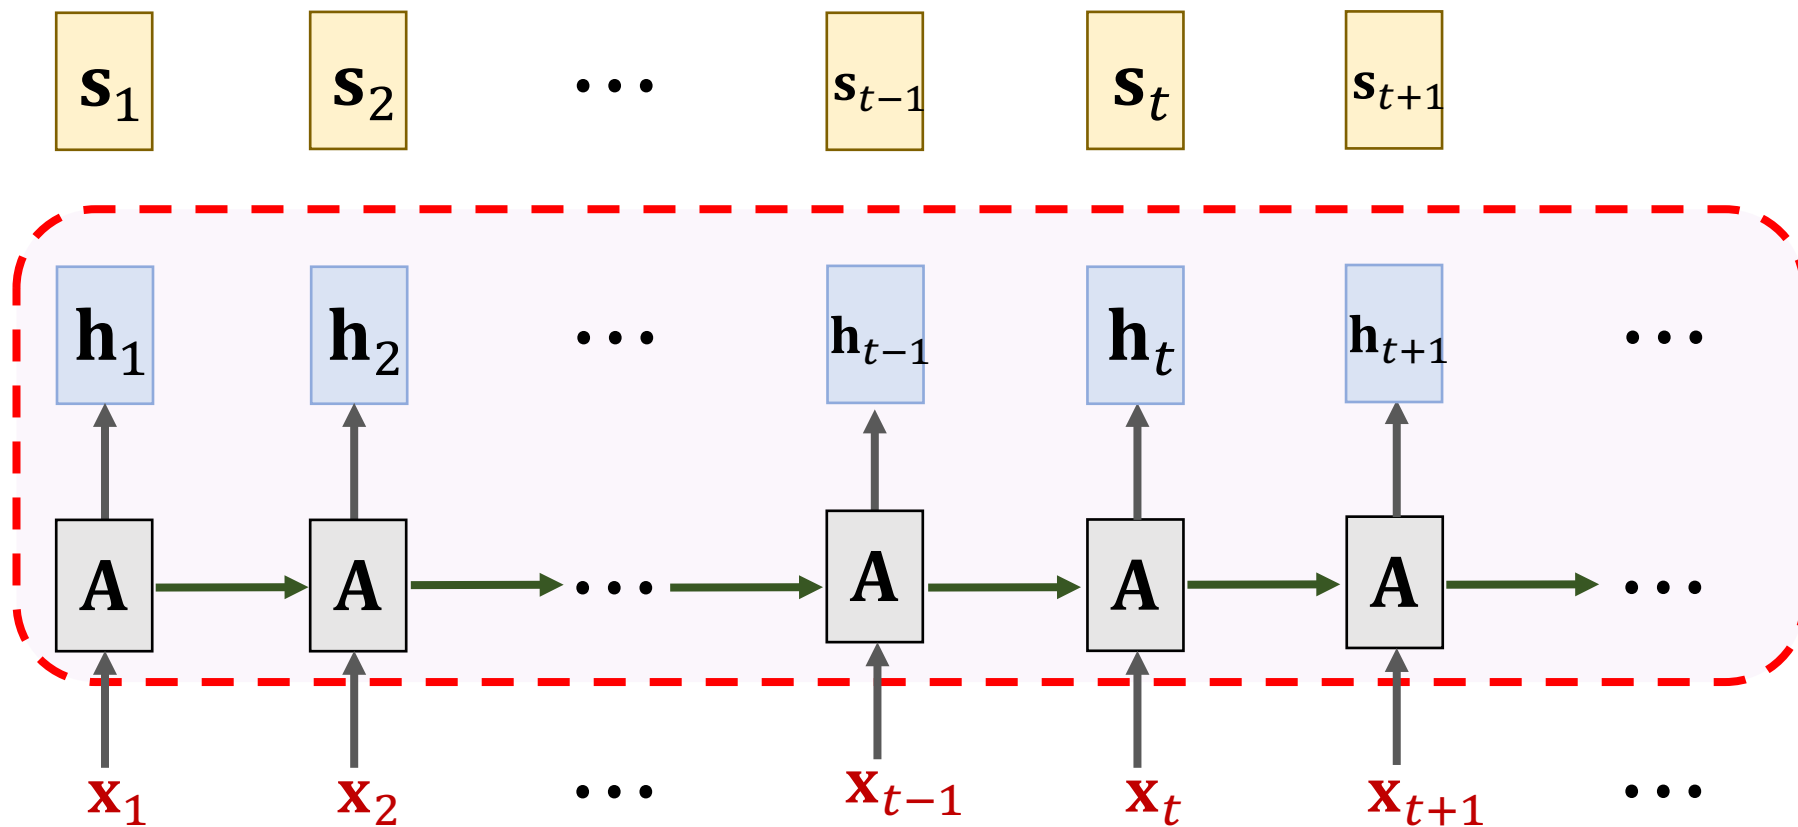


# SimpleRNN with Self-Attention

- $\mathbf{s}_{t+1} = \sum_{i=1}^t \alpha_i \mathbf{h}_i$ .
- $\alpha_i$  is computed by a neural network taking  $\mathbf{x}_{t+1}$ ,  $\mathbf{s}_t$ , and  $\mathbf{h}_i$  as inputs.




# SimpleRNN with Self-Attention






# LSTM with Self-Attention


- The update of LSTM with self-attention is analogous.



# LSTM with Self-Attention

- The update of LSTM with self-attention is analogous.
- $\mathbf{s}_t = \sum_{i=1}^{t-1} \alpha_i \mathbf{h}_i.$   Exactly the same as SimpleRNN

# LSTM with Self-Attention

- The update of LSTM with self-attention is analogous.
  - $\mathbf{s}_t = \sum_{i=1}^{t-1} \alpha_i \mathbf{h}_i.$
  - $\mathbf{c}'_t = \sum_{i=1}^{t-1} \alpha_i \mathbf{c}_i.$
- Conveyor belt
- 
- Exactly the same as SimpleRNN

# LSTM with Self-Attention

- The update of LSTM with self-attention is analogous.

- $\mathbf{s}_t = \sum_{i=1}^{t-1} \alpha_i \mathbf{h}_i.$

- $\mathbf{c}'_t = \sum_{i=1}^{t-1} \alpha_i \mathbf{c}_i.$

- $\mathbf{c}_t = \mathbf{f}_t \circ \mathbf{c}'_t + \mathbf{i}_t \circ \tilde{\mathbf{c}}_t.$

Forget gate, computed  
using  $\mathbf{x}_t$  and  $\mathbf{s}_t$ .

New value, computed using  $\mathbf{x}_t$  and  $\mathbf{s}_t$ .

Input gate, computed  
using  $\mathbf{x}_t$  and  $\mathbf{s}_t$ .

# LSTM with Self-Attention

- The update of LSTM with self-attention is analogous.
- $\mathbf{s}_t = \sum_{i=1}^{t-1} \alpha_i \mathbf{h}_i.$
- $\mathbf{c}'_t = \sum_{i=1}^{t-1} \alpha_i \mathbf{c}_i.$
- $\mathbf{c}_t = \mathbf{f}_t \circ \mathbf{c}'_t + \mathbf{i}_t \circ \tilde{\mathbf{c}}_t.$
- $\mathbf{h}_t = \mathbf{o}_t \circ \mathbf{c}_t.$



Output gate, computed  
using  $\mathbf{x}_t$  and  $\mathbf{s}_t$ .

# LSTM with Self-Attention

- The update of LSTM with self-attention is analogous.
- $\mathbf{s}_t = \sum_{i=1}^{t-1} \alpha_i \mathbf{h}_i.$
- $\mathbf{c}'_t = \sum_{i=1}^{t-1} \alpha_i \mathbf{c}_i.$
- $\mathbf{c}_t = \mathbf{f}_t \circ \mathbf{c}'_t + \mathbf{i}_t \circ \tilde{\mathbf{c}}_t.$
- $\mathbf{h}_t = \mathbf{o}_t \circ \mathbf{c}_t.$

## Difference 1:

- In standard LSTM, the gates and new value are computed using  $\mathbf{x}_t$  and  $\mathbf{h}_{t-1}$ .
- With self-attention, the gates and new value are computed using  $\mathbf{x}_t$  and  $\mathbf{s}_t$ .

# LSTM with Self-Attention

- The update of LSTM with self-attention is analogous.
- $\mathbf{s}_t = \sum_{i=1}^{t-1} \alpha_i \mathbf{h}_i.$
- $\mathbf{c}'_t = \sum_{i=1}^{t-1} \alpha_i \mathbf{c}_i.$
- $\mathbf{c}_t = \mathbf{f}_t \circ \mathbf{c}'_t + \mathbf{i}_t \circ \tilde{\mathbf{c}}_t.$
- $\mathbf{h}_t = \mathbf{o}_t \circ \mathbf{c}_t.$

## Difference 2:

- In standard LSTM, apply forget gate to  $\mathbf{c}_{t-1}.$
- With self-attention, apply forget gate to  $\mathbf{c}'_t = \sum_{i=1}^{t-1} \alpha_i \mathbf{c}_i.$

# Summary

- With self-attention, RNN is less likely to forget.



# Summary

- With self-attention, RNN is less likely to forget.
- Pay attention to the context relevant to the new input.

The  
The FBI  
The FBI is  
The FBI is chasing  
The FBI is chasing a  
The FBI is chasing a criminal  
The FBI is chasing a criminal on  
The FBI is chasing a criminal on the  
The FBI is chasing a criminal on the run  
The FBI is chasing a criminal on the run .

Figure is from the paper “ Long Short-Term Memory-Networks for Machine Reading.”